

Code Execution with MCP

Building More Efficient AI Agents

Based on the Anthropic Engineering article (Adam Jones & Conor Kelly)

The Problem

As MCP scales, agents drown in tokens

Why Agents Get Slow & Expensive

When agents connect to **hundreds or thousands of tools** across many MCP servers:


1. Tool definitions overload the context window

- All schemas loaded upfront
- Hundreds of thousands of tokens before the first user message

2. Intermediate results flow through the model

- Every tool output passes through context
- Large data gets re-written by the model into the next call

The Classic MCP Loop



Each tool call and result passes back through the model.

Example: Token Bloat

Task: "Download a 2-hour meeting transcript from Drive and attach it to a Salesforce le

```
TOOL CALL: gdrive.getDocument(...)  
  → 50,000-token transcript loaded into context  
  
TOOL CALL: salesforce.updateRecord(  
  data: { Notes: "...full transcript re-written..." }  
)  
  → model writes the entire transcript again
```

The transcript flows through context **twice**. Larger docs can exceed the context window

The Idea

Present MCP servers as a **code API**

Let the agent **write code** to call them

Tools as Files on a Filesystem

```
servers/  
├── google-drive/  
│   ├── getDocument.ts  
│   └── index.ts  
├── salesforce/  
│   ├── updateRecord.ts  
│   └── index.ts  
└── ...
```

Each tool = one file. The agent **explores the filesystem** and loads only the definitions it

Same Task — As Code

```
import * as gdrive from './servers/google-drive';
import * as salesforce from './servers/salesforce';

const transcript =
  (await gdrive.getDocument({ documentId: 'abc123' })).content;

await salesforce.updateRecord({
  objectType: 'SalesMeeting',
  recordId: '00Q5f000001abcXYZ',
  data: { Notes: transcript }
});
```

Result: ~150,000 tokens → ~2,000 tokens

A 98.7% reduction.

Five Big Wins

1. Progressive Disclosure

- Models are great at navigating filesystems
- Read tool definitions **on demand**
- Or expose a `search_tools` tool with a `detail_level` (name / name+desc / full schema)

Only pay tokens for tools you actually use.

2. Context-Efficient Results

Filter & transform **in the sandbox**, not in the model:

```
const allRows = await gdrive.getSheet({ sheetId: 'abc123' });  
const pending = allRows.filter(r => r.Status === 'pending');  
console.log(`Found ${pending.length} pending orders`);  
console.log(pending.slice(0, 5));
```

Model sees 5 rows. Not 10,000.

3. Real Control Flow

Loops, conditionals, retries — native code, not chained tool calls:

```
let found = false;
while (!found) {
  const msgs = await slack.getChannelHistory({ channel: 'C123' });
  found = msgs.some(m => m.text.includes('deployment complete'));
  if (!found) await new Promise(r => setTimeout(r, 5000));
}
```

Lower latency. No model round-trip per iteration.

4. Privacy-Preserving Operations

Intermediate results stay in the sandbox by default.

The MCP client can **tokenize PII** before the model ever sees it:

```
[  
  { salesforceId: '00Q...', email: '[EMAIL_1]',  
    phone: '[PHONE_1]', name: '[NAME_1]' },  
  ...  
]
```

Real values flow Drive → Salesforce. Model only sees tokens.

5. State Persistence & Skills

Agents can write to disk and resume later:

```
await fs.writeFile('./workspace/leads.csv', csvData);
```

And persist their **own code** as reusable skills:

```
// ./skills/save-sheet-as-csv.ts  
export async function saveSheetAsCsv(sheetId: string) { ... }
```

Over time the agent builds its own toolbox.

Trade-offs

Running agent-generated code requires:

Sandboxing · Resource limits · Monitoring

Weigh token & latency savings against the operational + security over

Takeaway

The novel problems of agents
have **classic software-engineering** solutions.

Let agents *write code* to talk to MCP.